CIALDINI INSTITUTE

Growth Agent Implementation Plan

Comprehensive Operational Memo for Execution

TO: Vasu (Senior Analyst, Principal Executor)

FROM: Samir Patel, CEO

DATE: March 21, 2026

RE: AI Growth Agent — Build, Deploy, and Operate

CLASSIFICATION: INTERNAL — CONFIDENTIAL

INTEGRITY · DISCIPLINE · SERVICE

1. Executive Summary

Vasu — this memo is your operating manual. We have built a working prototype of an AI-powered Growth Agent for Cialdini Institute. It generates scored growth ideas with execution plans, validates them through AI analysis, tracks them through a pipeline, and measures ROI. Your job is to take this from prototype to a production system that runs continuously and drives measurable audience growth, community memberships, and AI Bob subscriptions.

This document covers everything you need: the system architecture, what we've already built, what needs to happen next, the resources and budget required, the timeline, what I need to provide you, and exactly how we'll measure success.

Core thesis: We are execution-constrained, not idea-constrained. This system removes the bottleneck by generating, prioritizing, validating, and tracking growth initiatives at scale — so we can run 20–30 experiments simultaneously instead of agonizing over which one to start.

2. Strategic Context & Objectives

Everything we build must serve three measurable business outcomes, in priority order:

#	Strategic Priority	Primary KPI	12-Month Target
P1	Grow Cialdini Institute audience and community membership	Active community members	**5,000 active members**
P2	Drive AI Bob subscriptions	Paid subscribers	**1,000 paid subs**
P3	Establish CIP as the industry-standard influence assessment	CIP completions	**10,000 assessments taken**

The Growth Agent exists to systematically feed all three priorities with validated, tracked experiments. Our audience is global and virtual — we will target content and campaigns in English, Spanish, Portuguese, French, German, Arabic, Hindi, Mandarin, Japanese, and Korean from Day 1.

3. What We've Already Built (Prototype)

The Growth Agent prototype is a React-based dashboard application with persistent storage and live AI integration. Here is what currently works:

Feature	Current Capability	Status

Idea Generation	AI generates batches of 10–75 ideas with execution steps, cost estimates, validation methods, success metrics, target languages, and Cialdini principle mapping	**Working**
Scoring Engine	Automatic priority scoring (Impact × Ease matrix, 1–9 scale)	**Working**
AI Validation	One-click AI evaluation: proceed/kill/pivot verdict, risk assessment, quick-win alternatives, revised scoring	**Working**
Execution Planning	AI-generated week-by-week plans with tasks, milestones, tools needed, delegation notes, budget breakdown, go/no-go criteria	**Working**
Kanban Pipeline	6-stage visual pipeline: Idea → Validated → Executing → Measuring → Scaled / Killed	**Working**
ROI Tracker	Manual revenue/cost logging per idea, aggregate ROI dashboard	**Working**
Persistent Storage	All data persists across sessions	**Working**
Multi-Language Targeting	Ideas generated with language targeting across 10 languages	**Working**
Category Taxonomy	10 growth categories: Content/SEO, Community, Partnerships, Paid, Product, AI Bob, International, Events, Email, Referral	**Working**

4. System Architecture & Technical Specification

4.1 Current Stack

- Frontend: React (JSX) running as a Claude Artifact with persistent storage API
- AI Engine: Anthropic Claude Sonnet 4 via API (generation, validation, execution planning)
- Storage: Claude Artifact persistent key-value storage (JSON serialized)
- Deployment: Currently runs inside Claude.ai conversation interface

4.2 Production Architecture (Target State)

To make this a real operating system, we need to migrate from the prototype to a standalone deployment. Recommended architecture:

Layer	Technology	Rationale
Frontend	Next.js (React) deployed on Vercel or Cloudflare Pages	Fast, serverless, you already use Next.js for the investor portal
Backend API	Next.js API routes or Cloudflare Workers	Serverless, low-cost, handles AI calls securely
Database	Supabase (PostgreSQL) or PlanetScale	Structured data, real-time updates, familiar from investor portal work
AI Layer	Anthropic Claude API (Sonnet for speed, Opus for complex validation)	Best reasoning quality; we already have API access
Auth	Cloudflare Access or Supabase Auth	Role-gated: Samir (full), Vasu (full), Courtney (view + notes)
Scheduling	Cron jobs via Vercel or Cloudflare Workers	Automated weekly idea generation, daily metric pulls
Analytics Integration	Google Analytics, Squarespace Analytics, Stripe APIs	Pull real conversion data into ROI tracker automatically
Notifications	Slack webhooks or email (SendGrid)	Alert on validation results, milestone completions, ROI thresholds

4.3 AI Agent Capabilities Roadmap

The system should evolve through three capability tiers:

Tier 1 — Advisor (Current Prototype)

- Generates ideas on demand
- Validates ideas when asked
- Produces execution plans when triggered
- User manually moves ideas through pipeline stages
- User manually logs ROI data

Tier 2 — Operator (Target: Month 2–3)

- Scheduled weekly generation of 10–15 fresh ideas, auto-scored and categorized
- Auto-validates all new ideas within 24 hours of generation
- Sends Slack/email digest: “This week’s top 5 ideas to execute”
- Pulls real metrics from GA, Squarespace, Stripe to auto-populate ROI fields
- Courtney receives formatted task lists for delegation-ready execution steps

Tier 3 — Executor (Target: Month 4–6)

- Agent autonomously drafts blog posts, social copy, email sequences for approved ideas
- Agent creates localized content in target languages via translation + cultural adaptation

- Agent schedules social media posts, email sends through integrations
- Agent monitors metrics and auto-kills underperforming experiments after defined thresholds
- Agent suggests budget reallocation from killed experiments to scaled winners

5. Implementation Timeline

This is a 12-week build broken into 4 phases. Each phase has a clear deliverable and go/no-go gate.

PHASE 1: Foundation (Weeks 1–2)

Objective: Stand up the production environment and migrate the prototype.

Wk	Task	Owner
1	Set up Next.js project, Supabase database schema (ideas, notes, ROI, users), Cloudflare Access auth	Vasu
1	Migrate React prototype components into Next.js app	Vasu
1	Configure Anthropic API integration in backend (secure key handling, rate limiting)	Vasu
2	Build admin dashboard with Samir/Vasu/Courtney role gates	Vasu
2	Implement persistent storage (replace Artifact storage with Supabase)	Vasu
2	Deploy to Vercel/Cloudflare, internal testing	Vasu

Gate: *App is live behind auth, Samir/Vasu/Courtney can log in and generate + validate ideas.*

PHASE 2: Automation Layer (Weeks 3–5)

Objective: Move from manual triggers to scheduled, autonomous operations.

Wk	Task	Owner
3	Build cron job: weekly auto-generation of 15 ideas (Monday 6 AM ET)	Vasu
3	Build cron job: auto-validation of all new ideas within 24 hours	Vasu
3	Implement Slack webhook: weekly digest to #growth-ops channel	Vasu
4	Integrate Google Analytics API: pull pageviews, signups, conversions	Vasu

4	Integrate Squarespace Analytics: community page traffic, CIP completions	Vasu
4	Integrate Stripe API: AI Bob subscription data, revenue tracking	Vasu
5	Build Courtney's view: filtered task lists with execution steps, formatted for copy-paste	Vasu
5	Email notification system: daily summary of ideas in "Executing" stage	Vasu
5	End-to-end testing of full automation pipeline	Vasu + Samir

Gate: *System generates and validates ideas weekly without human trigger. Real metric data flows into ROI tracker. Courtney receives formatted task lists.*

PHASE 3: Content & Execution Engine (Weeks 6–9)

Objective: The agent starts producing real marketing assets, not just plans.

Wk	Task	Owner
6	Build content generation module: blog posts, LinkedIn posts, email sequences	Vasu
6	Build multi-language content pipeline: AI translation + cultural adaptation review	Vasu
7	Integrate social media scheduling (Buffer or native APIs)	Vasu
7	Build email draft integration (connect to existing email platform)	Vasu
8	Implement A/B test framework: agent creates 2–3 variants per content piece	Vasu
8	Build CIP promotion module: automated outreach sequences for assessment distribution	Vasu
9	Implement auto-kill logic: ideas below threshold after 14 days get flagged for review	Vasu
9	Build budget reallocation suggestions: agent recommends moving spend from killed to scaled	Vasu

Gate: *Agent produces publish-ready content in 3+ languages. At least 5 content pieces published through the system. A/B testing operational.*

PHASE 4: Optimization & Scale (Weeks 10–12)

Objective: Tighten feedback loops, optimize what's working, document the system.

Wk	Task	Owner

10	Build performance dashboard: real-time view of all experiments, aggregate ROI, category performance	Vasu
10	Implement learning loop: agent analyzes which categories/efforts/languages produce best ROI	Vasu
11	Build monthly executive report generator: automated PDF/slide deck of growth metrics	Vasu
11	Optimize AI prompts based on 10 weeks of data quality observations	Vasu + Samir
12	Full system documentation and runbook	Vasu
12	Training session for Courtney on daily operations	Vasu
12	Final review and formal handoff to operational mode	Samir + Vasu

Gate: *System is fully autonomous at Tier 2+ capability. Documentation complete. Courtney trained. Monthly reporting automated.*

6. Resources & Budget

6.1 Human Resources

Role	Person	Time Commitment	Duration
Principal Engineer	Vasu	30–40 hrs/week	12 weeks (Phases 1–4)
Executive Sponsor	Samir	3–5 hrs/week	Ongoing (reviews, inputs, approvals)
Admin / Ops	Courtney	5–10 hrs/week	Starting Phase 2 (Week 3+)
AI Prompt Tuning	Samir + Vasu	2–3 hrs/week	Weeks 1–6 (front-loaded)

6.2 Technology Costs (Monthly Estimates)

Service	Tier	Monthly Cost	Notes
Anthropic API (Claude)	Usage-based	$150–$400	~500K–1M tokens/week at Tier 2
Vercel or Cloudflare Pages	Pro	$20	Hosting + serverless functions
Supabase	Pro	$25	Database + auth + realtime

Cloudflare Access	Free–$5/user	$0–$15	3 users
SendGrid or Resend	Free tier	$0	Notifications + email
Buffer (social scheduling)	Essentials	$15	Optional; can use native APIs
Google Analytics	Free	$0	Already in use
Stripe API	Included	$0	Already paying Stripe fees

Estimated total monthly infrastructure cost: $210–$455/month. This does not include Vasu's or Courtney's time. The single largest variable cost is the Anthropic API, which scales with how aggressively we run the generation and validation engines.

7. Required Inputs from Samir

Vasu, you will need the following from me before or during each phase. I'm listing them here so Courtney can help schedule these and so we don't create blockers.

7.1 Before Phase 1 Begins (Pre-Work)

1. **Anthropic API key** with appropriate rate limits and billing configured. Confirm whether we use the existing key or provision a dedicated one for the Growth Agent.
2. **Deployment target decision:** Vercel (simpler, faster) vs. Cloudflare Pages (consistent with existing infra). Need a yes/no.
3. **Supabase project credentials** or authorization for Vasu to create a new project under the existing org.
4. **Domain/subdomain:** e.g., growth.cialdini.com or ops.patel-capital.net. Need DNS access or a decision.
5. **Role permissions matrix:** Confirm what Courtney can and cannot do (generate? validate? delete? publish content?).

7.2 During Phase 2 (Automation)

1. **Google Analytics property ID** and service account credentials for API access.
2. **Squarespace API key** or confirmation of which analytics endpoints we can access.
3. **Stripe API key (read-only)** for pulling AI Bob subscription data and revenue.
4. **Slack workspace:** Create a #growth-ops channel and provide webhook URL, or authorize Vasu to set this up.
5. **Weekly 30-minute review:** Commit to a standing meeting (suggest Mondays) to review the week's generated ideas and approve/reject validation results.

7.3 During Phase 3 (Content Engine)

1. **Brand voice document:** A 1–2 page brief on Cialdini Institute's tone, voice, and content guidelines. If this doesn't exist, we need 30 minutes to create one together.
2. **Content approval workflow:** Decision on whether AI-generated content needs Samir's approval before publishing, or if Vasu/Courtney can approve within guidelines.
3. **Social media credentials:** Access to LinkedIn, Twitter/X, and any other platforms for scheduling integration.
4. **Email platform access:** Credentials for the email marketing platform (Mailchimp, ConvertKit, etc.) used for Cialdini Institute communications.
5. **Translation quality review contacts:** Native speakers for our top 3 non-English markets (likely Spanish, Portuguese, Mandarin) who can spot-check AI translations before launch.

7.4 Ongoing (All Phases)

1. **Strategic guardrails:** Any categories, tactics, or markets that are off-limits (e.g., no paid ads on certain platforms, no cold outreach to certain audiences, compliance constraints for Reg D).
2. **Monthly budget cap:** Maximum monthly spend on paid experiments generated by the system. Recommend starting at $500/month and scaling based on ROI.
3. **Kill criteria:** At what point do we kill an experiment? Recommend: no measurable traction after 14 days for low-effort, 30 days for high-effort.

8. Success Metrics & Reporting

We will track the Growth Agent's own performance, separate from the business KPIs it serves. This is how we know if the system itself is working.

System Metric	Week 4 Target	Week 8 Target	Week 12 Target
Total ideas generated	100+	300+	500+
Ideas validated by AI	80+	250+	450+
Ideas in "Executing" or beyond	10+	30+	50+
Ideas with ROI data logged	5+	20+	40+
Positive-ROI experiments	2+	10+	20+
Content pieces published via system	—	10+	40+

Languages with active content	1 (EN)	3+	5+
Time from idea to execution	3–5 days	1–2 days	< 24 hours (auto)

The monthly executive report (automated in Phase 4) will include aggregate ROI by category, top-performing and worst-performing experiments, language/market breakdown, and recommendations from the AI for the next month's focus areas.

9. Risk Register

Risk	Likelihood	Impact	Mitigation
AI generates low-quality or generic ideas	Medium	Wasted execution time on bad ideas	Validation gate catches these; tune prompts weekly with real examples of good vs. bad
API costs spike unexpectedly	Low	Budget overrun	Set hard API spend limits in Anthropic dashboard; use Sonnet (cheaper) for generation, Opus only for complex validation
Multi-language content quality issues	Medium	Brand damage in international markets	Human spot-check for top 3 languages before publishing; start with English + Spanish only, expand gradually
Integration failures (GA, Stripe, etc.)	Medium	ROI tracker goes blind	Build manual fallback entry; decouple metric ingestion from core pipeline so failures don't block operations
Vasu bandwidth constraints	Medium	Timeline slippage	Front-load Phase 1 when it's pure build work; Phase 2–4 overlap with operational duties is lighter
Reg D compliance: marketing to non-accredited investors	Low–Med	Regulatory risk for Trophy Point cross-promotion	Growth Agent does NOT generate Trophy Point investor acquisition content; that stays separate and compliance-reviewed

10. Immediate Next Steps

Within the next 48 hours, the following should happen:

#	Action	Owner	Deadline
1	Vasu reviews this memo and flags questions or concerns	Vasu	**Day 1**
2	Samir provides API key, deployment decision, and domain choice (Section 7.1 items 1–4)	Samir	**Day 2**
3	Samir confirms Courtney's permission level (Section 7.1 item 5)	Samir	**Day 2**
4	Vasu sets up the GitHub repo, Next.js project scaffold, and Supabase project	Vasu	**Day 3**
5	Courtney schedules standing Monday review meeting (30 min, Samir + Vasu)	Courtney	**Day 3**
6	Vasu begins Phase 1 build	Vasu	**Day 4**

11. Closing Note

Vasu — I built the prototype because I was scared of executing. **This system exists to make execution mechanical, not emotional.** The AI generates the ideas. The AI validates them. The AI writes the plan. Your job is to build the machine that lets the AI do that continuously, and Courtney's job is to push the buttons it tells her to push.

We are not trying to be perfect. We are trying to run 30 experiments at once and let the data tell us which 3 to scale. Speed and volume beat perfection here.

Let's build this. — Samir

Samir Patel
Chief Executive Officer, Cialdini Institute LLC

CEO, Trophy Point Investment Group